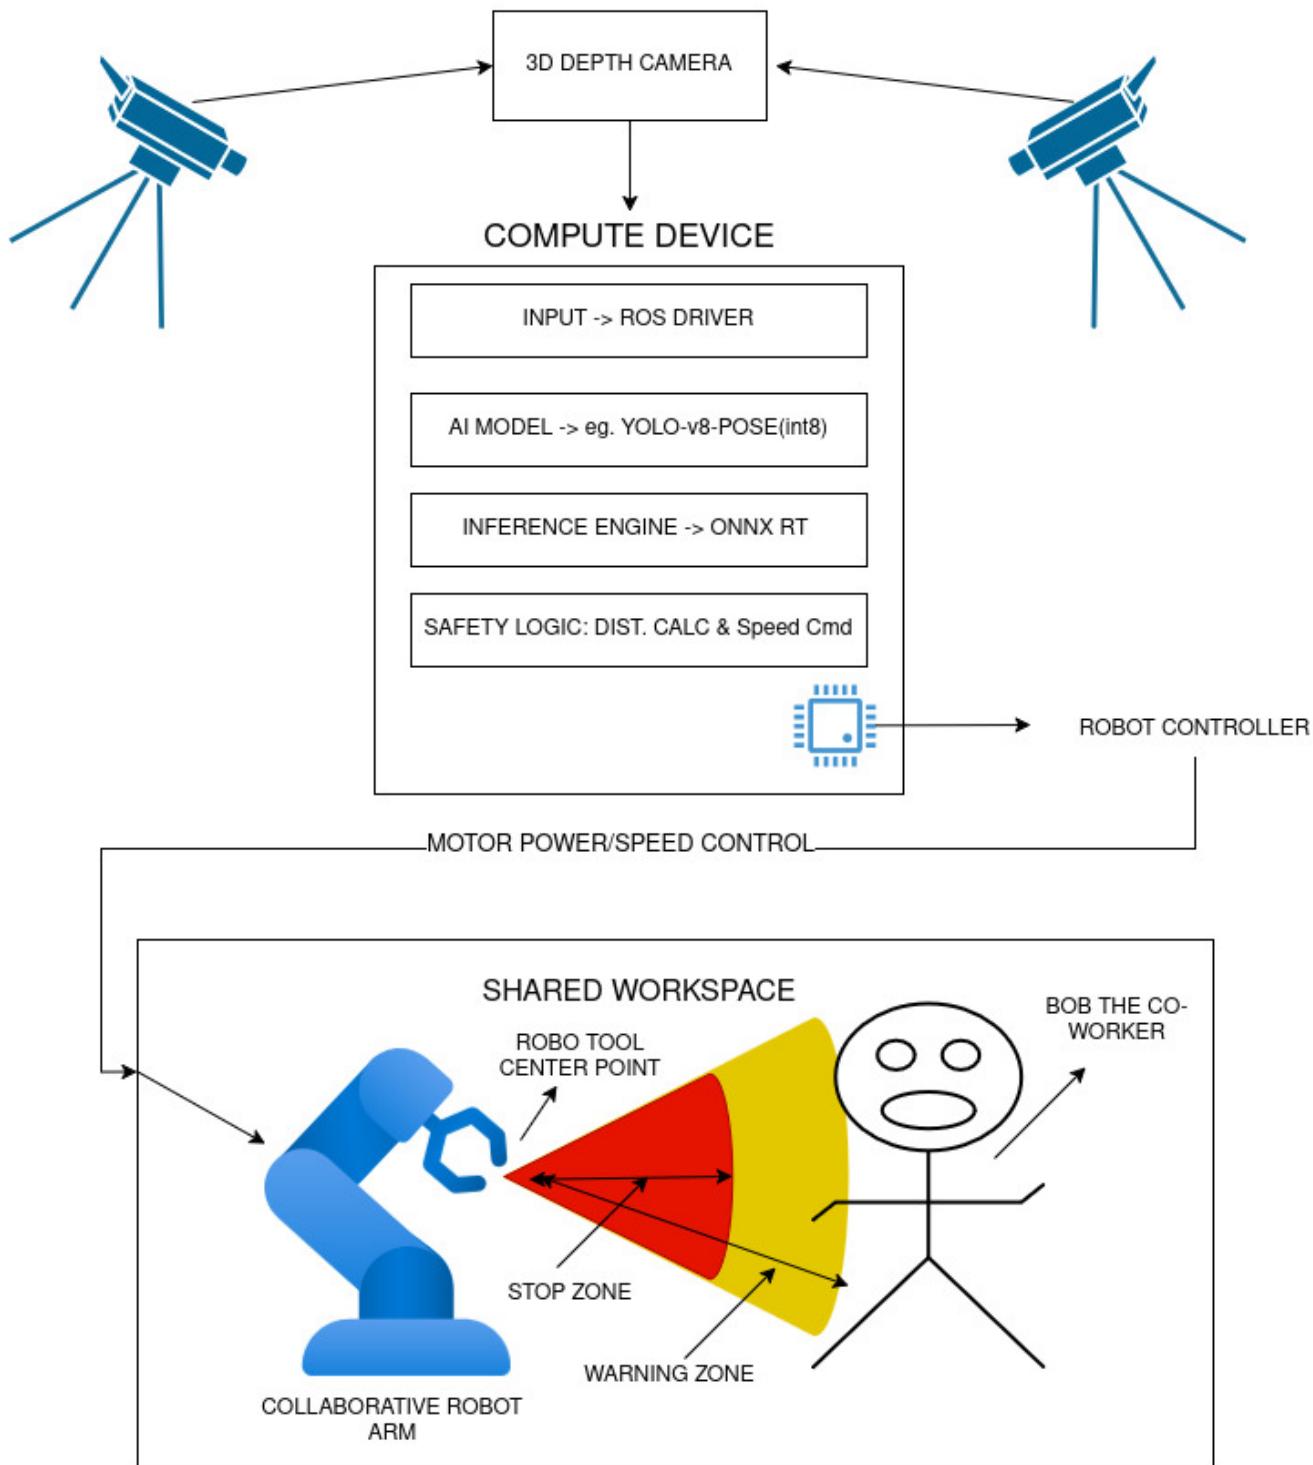


Edge AI Architecture for Human-Robot Collaboration

Transitioning from Rigid Safety Fencing to Dynamic Edge AI Safety System.

Speed and Separation Monitoring (SSM)



LOW LATENCY EDGE PROCESSING FOR REAL_TIME SAFETY

1. Summary

Traditional industrial safety relies on physical cages that completely stop production when a human enters the workspace. This "Stop-and-Go" approach kills efficiency. Modern Industry 4.0 standards (ISO/TS 15066) allow for **Speed and Separation Monitoring (SSM)**, where the robot dynamically slows down or alters its path based on the human's proximity.

However, relying on Cloud AI for this task introduces dangerous latency. A 100ms network lag could result in an injury before the robot receives the stop command. This report outlines a **High-Performance Edge AI Architecture** that processes skeletal tracking locally (within 15ms), ensuring compliance with safety standards while maintaining high production throughput.

2. System Architecture Design

The proposed system creates a "Virtual Safety Shield" around the collaborative robot (Cobot). It uses redundant 3D vision to track the human worker's skeletal keypoints (wrists, head, shoulders) in real-time.

2.1 Hardware Components

Component	Specification	Function
Vision Sensors	2x Intel RealSense D435i	Redundancy & Depth: Two cameras positioned at opposing angles eliminate "blind spots" (occlusion) caused by the robot's own arm. These provide both RGB and Depth (Point Cloud) data.
Edge Compute	NVIDIA Jetson Nano/Orin	Local Inference: A dedicated System-on-Module (SoM) It processes video streams locally without sending data to the cloud.
Robot Controller	Universal Robots Control Box	Motion Control: Receives speed scaling commands (0% to 100%) from the Edge device and executes Safety Stop Category 1 or 2.

2.2 Software Stack

- **Input Layer:** ROS 2 (Robot Operating System) nodes capturing raw point cloud data at high fps.
- **AI Model: YOLOv8-Pose**
 - Pose estimation tracks specific joints (e.g., the right wrist). This allows the robot to keep working if the human is nearby but facing away, only stopping if the *hand* reaches toward the danger zone.
- **Inference Engine: ONNX RT**
 - This is the critical optimization layer that compiles the generic PyTorch model into a binary engine optimized.
- **Safety Logic:**
 - A deterministic algorithm that calculates the Euclidean distance between the robot's **Tool Center Point (TCP)** and the nearest **Human Keypoint**.

3. Edge AI Optimization Strategy

Running a modern computer vision model like YOLOv8-Pose at 30+ FPS on a small embedded device requires aggressive optimization. We cannot use the "heavy" models used in cloud servers.

Strategy A: Post-Training Quantization to INT8

Standard AI models compute using **FP32** (32-bit Floating Point) numbers. This is highly precise but computationally expensive.

- **The Technique:** We convert the model weights and activations to **INT8** (8-bit Integers).
- **Implementation:** Using the TensorRT calibrator, we pass a small dataset of "representative images" through the model. The calibrator determines the dynamic range of activation values and maps them to the -128 to +127 integer range.
- **Benefit:** This reduces memory bandwidth usage by **4x** and inference time is lot faster.

Strategy B: Structured Pruning

Many neural networks are "over-parameterized," containing neurons that contribute little to the final output.

- **The Technique:** We use **Structured Pruning** (Channel Pruning). Unlike *unstructured* pruning (which makes matrices sparse but doesn't necessarily speed up hardware), structured pruning removes entire filters (channels) from the Convolutional layers.
- **Implementation:** We identify channels with the lowest L1-norm (least impact) and physically remove them from the network architecture, effectively creating a "thinner" model.
- **Benefit:** This directly reduces the number of FLOPs (Floating Point Operations) required for inference, lowering latency.

4. Impact Analysis: Cloud vs. Edge

According to ISO 13855, the Minimum Safety Distance (S) is calculated as:

$$S = (K \times T) + C$$

- K : Human approach speed (1600 mm/s).
- T : Total system stopping time (Latency + Braking time).

If Latency (T) increases, the Safety Distance (S) must increase, forcing the robot to stop when the human is still far away, ruining productivity.

Note: metrics are relative to each other

Metric	Cloud Solution (Standard)	Edge AI Solution (Optimized)	Impact Analysis
Total Latency (T)	HIGH (Network RTT + Inference)	LOW (Local Capture + Inference)	Reducing T allows the robot to work much closer to the human without violating safety zones.
Energy	HIGH (server side)	LOW	Compute devices like Jetson runs on low voltage, easily integrated into the robot's power supply without expensive cooling.

Metric	Cloud Solution (Standard)	Edge AI Solution (Optimized)	Impact Analysis
Model Accuracy	HIGH	LOW	The precision loss from INT8 quantization is negligible for detecting a large object like a human arm at close range.

5. Risk Assessment:

While optimization is necessary, aggressive compression introduces specific safety risks that must be mitigated.

Risk 1: "The Ghost Effect" (False Negatives)

Description: Aggressive **Pruning** removes "redundant" features. In rare edge cases—such as a worker wearing a grey shirt lying on a grey floor—the compressed model might lack the feature extractors necessary to distinguish the arm from the background.

- **Consequence:** The system fails to detect the human (False Negative), and the robot does not stop.

Risk 2: "The Jitter Effect" (Temporal Instability)

Description: **Quantization** (INT8) introduces rounding errors. While accurate on average, the detected coordinate of a wrist might jump by 5-10cm between frames even if the human is still.

- **Consequence:** The robot controller receives erratic distance data, causing the robot to "stutter" (rapidly braking and accelerating).

6. Conclusion

By deploying a **Quantized and Pruned YOLOv8-Pose model** on an **Edge device**, we can achieve a safety system that is compliant with **ISO/TS 15066**. This architecture solves the latency issues of cloud computing, enabling an environment where humans and robots work in close proximity without physical barriers.